

CSC420 Assignment 2

1. Implement seam carving

a) Compute magnitude of gradients of an image.

```

6   def rgb_gradient(img):
7       """
8           Compute magnitude of gradients for the input RGB image.
9       """
10      height, width = img.shape[0], img.shape[1]
11
12      # rgb channels
13      red_channel = np.zeros((height, width))
14      green_channel = np.zeros((height, width))
15      blue_channel = np.zeros((height, width))
16
17      for row in range(height):
18          for col in range(width):
19              red_channel[row][col] = img[row][col][0]
20              green_channel[row][col] = img[row][col][1]
21              blue_channel[row][col] = img[row][col][2]
22
23      # compute x- and y- gradients for each channel
24      red_x = ndimage.sobel(red_channel, axis=0, mode='constant')
25      red_y = ndimage.sobel(red_channel, axis=1, mode='constant')
26      green_x = ndimage.sobel(green_channel, axis=0, mode='constant')
27      green_y = ndimage.sobel(green_channel, axis=1, mode='constant')
28      blue_x = ndimage.sobel(blue_channel, axis=0, mode='constant')
29      blue_y = ndimage.sobel(blue_channel, axis=1, mode='constant')
30
31      result = np.zeros((height, width))
32
33      # compute the magnitude of gradients for each pixel
34      for row in range(height):
35          for col in range(width):
36              red = np.hypot(red_x[row, col], red_y[row, col])
37              blue = np.hypot(blue_x[row, col], blue_y[row, col])
38              green = np.hypot(green_x[row, col], green_y[row, col])
39              result[row, col] = np.sqrt(red ** 2 + green ** 2 + blue ** 2)
40
41      return result

```

- b) Find the connected path of pixels that has the smallest sum of gradients.

```

44     def cost(gradients):
45         """
46             For each pixel, calculate the smallest sum of gradients of path that
47             starts from the pixel to the end row of the image.
48         """
49         row, col = gradients.shape
50         pixel_cost = np.zeros(gradients.shape)
51
52         # the bottom-most row, cost = the gradient of the pixel
53         pixel_cost[row - 1, :] = gradients[row - 1, :]
54
55         for i in range(row - 2, -1, -1):
56             for j in range(col):
57                 # only consider connected neighboring pixels
58                 # connect to one of the column j-1, j, j+1 with minimum cost
59                 left, right = max(j - 1, 0), min(col, j + 2)
60                 pixel_cost[i][j] = gradients[i][j] + pixel_cost[i + 1, left: right].min()
61
62         return pixel_cost
63
64
65     def find_seam_path(pixel_cost):
66         """
67             Find the connected path of pixels that has the smallest sum of gradients.
68         """
69         row, col = pixel_cost.shape
70
71         path = []
72         # append the index of the minimum cost among the first row
73         j = pixel_cost[0].argmin()
74         path.append(j)
75
76         for i in range(row - 1):
77             # only consider connected neighboring pixels
78             left, right = max(j - 1, 0), min(col, j + 2)
79             # start from the second row
80             j = max(j - 1, 0) + pixel_cost[i + 1, left: right].argmin()
81             path.append(j)
82
83     return path

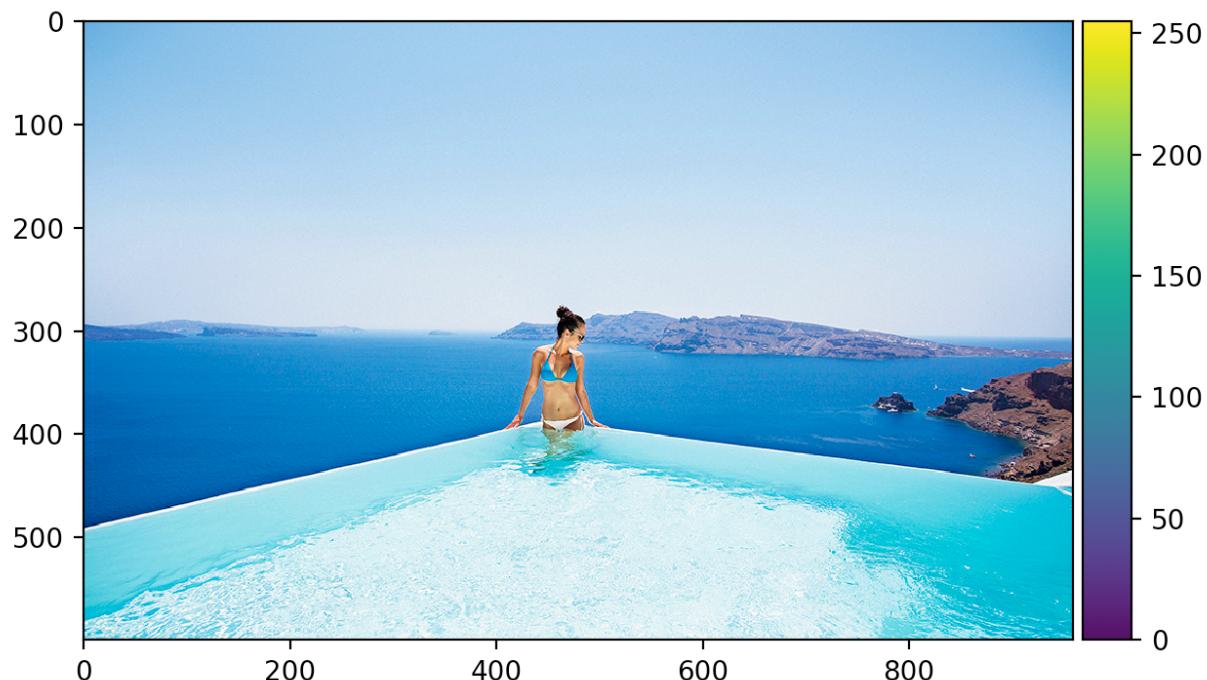
```

- c) Remove the pixels in the path from the image. Get a new image with one column less.

```

86     def remove_seam_path(img, path):
87         """
88             Remove the pixels in the path from the image.
89         """
90         row, col = img.shape[0], img.shape[1]
91         new_img = np.zeros(img.shape)
92
93         for i, j in enumerate(path):
94             # copy
95             new_img[i, 0:j, :] = img[i, 0:j, :]
96             # remove the jth column with minimum cost
97             new_img[i, j:col-1, :] = img[i, j+1:col, :]
98
99         res_img = new_img[:, :-1, :].astype(np.int64)
100        return res_img
101
102
103    def seam_carving(img):
104        """
105            Implement seam carving.
106        """
107        # (a) Compute magnitude of gradients of an image
108        gradients = rgb_gradient(img)
109
110        # (b) Find the connected path of pixels with the smallest sum of gradients.
111        pixel_cost = cost(gradients)
112        seam_path = find_seam_path(pixel_cost)
113
114        # (c) Remove the pixels in the path from the image.
115        new_img = remove_seam_path(img, seam_path)
116
117        return new_img

```

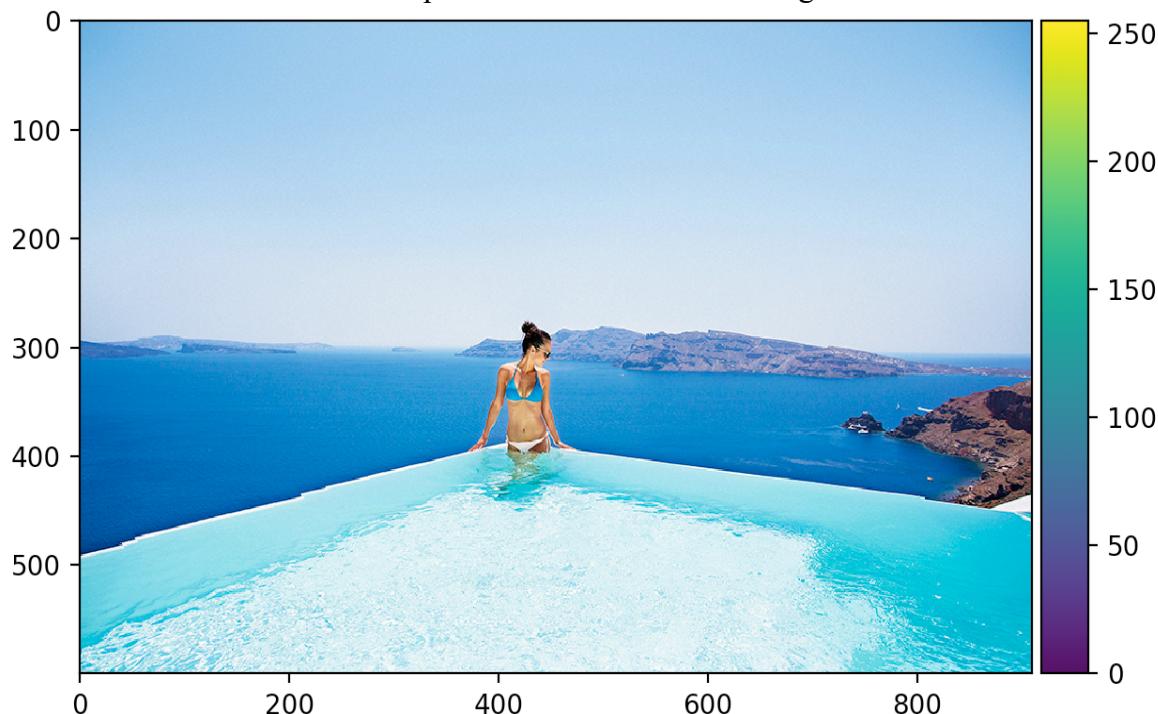


d) Remove a few paths with the lowest sum of gradients. Create a few examples.

i. Remove 10 paths with the lowest sum of gradients.



ii. Remove 50 paths with the lowest sum of gradients.

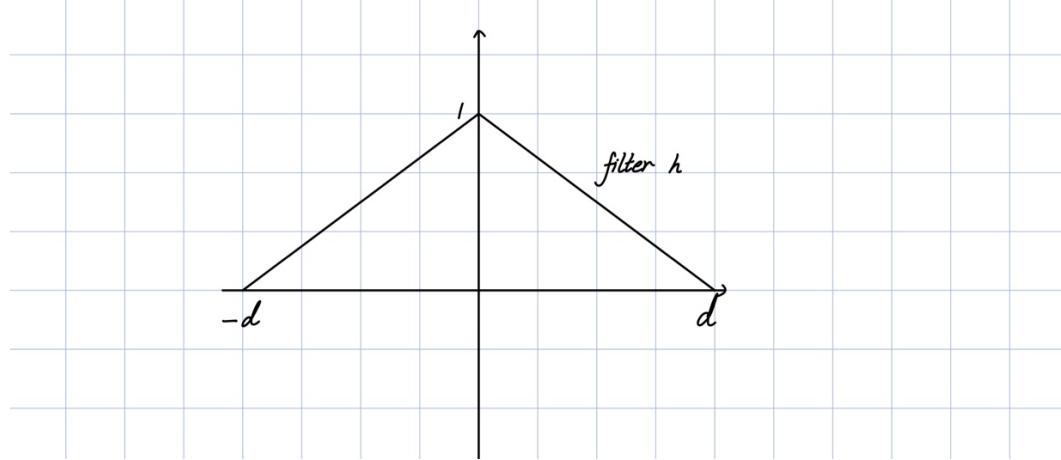


2. Image Upscaling

- 1) We could use the linear interpolation method and then the linear convolution filter h would have the mathematical form:

$$h = [0, 1/d, \dots, (d-1)/d, 1, (d-1)/d, \dots, 1/d, 0],$$
 where d is the upscaling factor.

- 2) Plot this filter.



3. Neural Network

```

10 def cross_entropy_loss_function(prediction, label):
11     # TODO: compute the cross entropy loss function between the prediction and ground truth label.
12     # prediction: the output of a neural network after softmax. It can be an Nxd matrix, where N is the
13     #             number of samples and d is the number of different categories
14     # label: The ground truth Labels, it can be a vector with length N, and each element in this vector
15     # Note: we take the average among N different samples to get the final loss.
16     # number of samples
17     n = prediction.shape[0]
18
19     # Compute the cross-entropy loss
20     log_probs = np.dot(np.log(prediction), label.T) + np.dot(np.log(1 - prediction), (1 - label).T)
21     loss = - np.sum(log_probs) / n
22
23     return loss
24
25
26 def sigmoid(x):
27     # TODO: compute the softmax with the input x: y = 1 / (1 + exp(-x))
28     return 1 / (1 + np.exp(-x))
29
30
31 def sigmoid_derivative(s):
32     return s * (1 - s)
33
34
35 def softmax(x):
36     # TODO: compute the softmax function with input x.
37     # Suppose x is Nxd matrix, and we do softmax across the Last dimension of it.
38     # For each row of this matrix, we compute x_{j, i} = exp(x_{j, i}) / \sum_{k=1}^d exp(x_{j, k})
39     return np.exp(x) / np.sum(np.exp(x), axis=0)
40

```

```

42     class OneLayerNN():
43         def __init__(self, num_input_unit, num_output_unit):
44             #TODO: Random Initliaize the weight matrixs for a one-layer MLP.
45             # the number of units in each layer is specified in the arguments
46             # Note: We recommend using np.random.randn() to initialize the weight ma
47             #         and initialize the bias matrix as full zero using np.zeros()
48
49             self.w1 = np.random.randn(num_output_unit, num_input_unit)
50             self.b1 = np.zeros((num_output_unit, 1))
51
52         def forward(self, input_x):
53             #TODO: Compute the output of this neural network with the given input.
54             # Suppose input_x is an Nxd matrix, where N is the number of samples and
55             # Compute output: z = softmax (input_x * W_1 + b_1), where W_1, b_1 are
56             # Note: If we only have one layer in the whole model and we want to use
57             #       then we directly apply softmax **without** using sigmoid (or rel
58
59             z1 = np.dot(self.w1, input_x.T) + self.b1
60             a1 = softmax(z1)
61
62             return a1
63
64         def backpropagation_with_gradient_descent(self, loss, learning_rate, input_x):
65             #TODO: given the computed loss (a scalar value), compute the gradient fr
66             # Note that you may need to store some intermediate value when you do fo
67             # Suggestions: you need to first write down the math for the gradient, t
68
69             n = input_x.shape[0]
70             a1 = self.forward(input_x)
71
72             dl = a1 - label
73             # the derivative of softmax here is softmax * (1 - softmax)
74             dz1 = np.dot(dl, (a1 * (1 - a1)).T)
75             dw1 = (1 / n) * np.dot(dz1, input_x.T)
76             db1 = (1 / n) * np.sum(dz1, axis=1, keepdims=True)
77
78             self.w1 = self.w1 - learning_rate * dw1
79             self.b1 = self.b1 - learning_rate * db1

```

[Additional credit]

```

82     # [Bonus points] This is not necessary for this assignment
83     class TwoLayerNN():
84         def __init__(self, num_input_unit, num_hidden_unit, num_output_unit):
85             #TODO: Random Initialize the weight matrixs for a two-layer MLP with sigmoid
86             # the number of units in each layer is specified in the arguments
87             # Note: We recommend using np.random.randn() to initialize the weight matrix
88             #       and initialize the bias matrix as full zero using np.zeros()
89
90             self.w1 = np.random.randn(num_hidden_unit, num_input_unit)
91             self.b1 = np.zeros((num_hidden_unit, 1))
92
93             self.w2 = np.random.randn(num_output_unit, num_hidden_unit)
94             self.b2 = np.zeros((num_output_unit, 1))
95
96         def forward(self, input_x):
97             #TODO: Compute the output of this neural network with the given input.
98             # Suppose input x is Nxd matrix, where N is the number of samples and d is the dimension
99             # Compute: first layer: z = sigmoid (input_x * W_1 + b_1) # W_1, b_1 are weight and bias
100            # Compute: second layer: o = softmax (z * W_2 + b_2) # W_2, b_2 are weight and bias
101
102             z1 = np.dot(self.w1, input_x.T) + self.b1
103             a1 = sigmoid(z1)
104
105             z2 = np.dot(self.w2, a1.T) + self.b2
106             a2 = softmax(z2)
107
108             return a2
109
110
111         def backpropagation_with_gradient_descent(self, loss, learning_rate, input_x, label):
112             #TODO: given the computed Loss (a scalar value), compute the gradient for the weights and biases
113             # Note that you may need to store some intermediate value when you do forward pass
114             # Suggestions: you need to first write down the math for the gradient, then implement it
115
116             n = input_x.shape[0]
117             z1 = np.dot(input_x, self.w1) + self.b1
118             a1 = sigmoid(z1)
119             a2 = self.forward(input_x)
120
121             dl = a2 - label
122             # softmax
123             dz2 = np.dot(dl, (a2 * (1 - a2)).T)
124             dw2 = (1 / n) * np.dot(dz2, a1.T)
125             db2 = (1 / n) * np.sum(dz2, axis=1, keepdims=True)
126             # sigmoid
127             dz1 = np.dot(dw2, sigmoid_derivative(z1).T)
128             dw1 = (1 / n) * np.dot(dz1, input_x.T)
129             db1 = (1 / n) * np.sum(dz1, axis=1, keepdims=True)
130
131             self.w1 = self.w1 - learning_rate * dw1
132             self.b1 = self.b1 - learning_rate * db1
133             self.w2 = self.w2 - learning_rate * dw2
134             self.b2 = self.b2 - learning_rate * db2

```