

CSC420 Assignment 4

1. Estimate the width and height of the door (in cm) from the picture.



Note: There is a manual taped on the door, with a length of 21.6cm (~816 pixels) and a width of 9.2cm (~348 pixels).

Output:

```
Run: door x
/Users/mac/Desktop/22Winter/CSC420/A4/Assignment4/venv/bin/python /
width of the door 197.0002251381004
height f the door 82.8642354586023
Process finished with exit code 0
```

Implementation

```

1  import numpy as np
2      import cv2
3      import matplotlib.pyplot as plt
4      from skimage import transform
5  from scipy.spatial.distance import euclidean
6
7  if __name__ == '__main__':
8      door_img = cv2.imread('door.JPG', 1)[: :, ::-1]
9
10     door_coordinates = np.array([[312, 473], [1681, 348],
11                                   [1468, 4008], [372, 3598]])
12     manual_coordinates = np.array([[1172, 1880], [1314, 1887],
13                                    [1302, 2273], [1159, 2246]])
14
15     x_manual = list(manual_coordinates[:, [0]].flatten())
16     y_manual = list(manual_coordinates[:, [1]].flatten())
17     x_door = list(door_coordinates[:, [0]].flatten())
18     y_door = list(door_coordinates[:, [1]].flatten())
19
20     # outline the door by green lines
21     plt.plot(x_door + x_door[:1], y_door + y_door[:1], 'g')
22     # outline the manual by red lines
23     plt.plot(x_manual + x_manual[:1], y_manual + y_manual[:1], 'r')
24
25     plt.imshow(door_img)
26     plt.show()
27
28     f = 38 # 1cm ~ 38 pixels
29     x_real_manual = [0 + 1, 9.2 * f + 1, 9.2 * f + 1, 0 + 1]
30     y_real_manual = [0 + 1, 0 + 1, 21.6 * f + 1, 21.6 * f + 1]
31
32     src = np.vstack((x_manual, y_manual)).T
33     dst = np.vstack((x_real_manual, y_real_manual)).T
34
35     # compute homography H
36     tform = transform.estimate_transform('projective', src, dst)
37     src_door = np.row_stack((x_door, y_door, [1] * 4))
38     x_homography = tform.params.dot(src_door)
39     x_homography = x_homography / x_homography[2, :]
40
41     # compute the coordinates of the corners
42     upper_right = [x_homography[0, :][np.argmin(x_homography[1, :])],
43                   x_homography[1, :][np.argmin(x_homography[1, :])]]
44
45     bottom_right = [x_homography[0, :][np.argmax(x_homography[1, :])],
46                     x_homography[1, :][np.argmax(x_homography[1, :])]]
47
48     upper_left = [x_homography[0, :][0], x_homography[1, :][0]]
49
50     # compute the distance and convert to cm unit
51     width = euclidean(upper_left, upper_right) / f
52
53     height = euclidean(upper_right, bottom_right) / f
54
55     print('width of the door {}'.format(width))
56     print('height of the door {}'.format(height))

```

2. Take two photos, landscape 1 and landscape 2 and stitch them into one photograph.

a. Stitched output image:



b. Implementation

```

1  from SIFT_matching import * # Implemented in A3
2  import numpy as np
3  import random
4
5
6  def get_matching_keypoints(img1, img2):
7      """
8      Return a list of matched pair tuples (keypoint_1, keypoint_2) for
9      two input images.
10     """
11     keypoint_1, des_1 = extract_SIFT_keypoints(img1)
12     keypoint_2, des_2 = extract_SIFT_keypoints(img2)
13
14     matched_keypoints = []
15
16     for i in range(len(keypoint_1)):
17         # match keypoint corresponding to each keypoint
18         f_j, j, thresh_ratio = calculate_correspondance(des_1[i], des_2)
19
20         # if ratio between closet and second closet match is < 0.8, keep match
21         if thresh_ratio < 0.8:
22             match_tuple = (keypoint_1[i], keypoint_2[j])
23             matched_keypoints.append(match_tuple)
24
25     return matched_keypoints
26
27
28  def calculate_homography(matched_keypoints):
29      """
30      Given a list of matched pair tuples (keypoint_1, keypoint_2) for both images,
31      return the Homography matrix (H)
32      """
33
34     length = len(matched_keypoints)
35     A = np.empty((0, 9))
36
37     for i in range(length):
38         (keypoint_i, keypoint_j) = matched_keypoints[i]
39         (x_j, y_j) = keypoint_i.pt
40         (x_k, y_k) = keypoint_j.pt

```

```

41
42     # append partial to A
43     a = np.array([[x_j, y_j, 1, 0, 0, 0, -x_k*x_j, -x_k*y_j, -x_k],
44                   [0, 0, 0, x_j, y_j, 1, -y_k*x_j, -y_k*y_j, -y_k]])
45     A = np.append(A, a, axis=0)
46
47     # eigenvalues and eigenvectors
48     w, v = np.linalg.eig(A.T @ A)
49
50     # homography H
51     i_min = np.argmin(w)
52     h = v[:, i_min]
53     H = np.reshape(h, (3,3))
54
55     return H
56
57
58 def transform_point(H, point):
59     """
60     Helper to transform points.
61     """
62     (x, y) = point
63     q = np.array([x, y, 1])
64     wp = H @ q
65     p = wp/wp[2]
66
67     return p[0], p[1]
68
69
70 def count_inliers(H, matched_keypoints):
71     """
72     Helper to compute the number of inliers.
73     """
74     n = 0
75     for i in range(len(matched_keypoints)):
76         (keypoint_i, keypoint_j) = matched_keypoints[i]
77         p = np.asarray(keypoint_j.pt)
78         p_hat = np.asarray(transform_point(H, keypoint_i.pt))
79         d = np.linalg.norm(p - p_hat)
80
81         if d < 3:
82             n += 1
83

```

```

84     return n
85
86
87 def RANSAC_homography(matched_keypoints):
88     """
89     Given a list of matched keypoints, return the best solution for the
90     homography (H) calculated via RANSAC.
91     """
92
93     # iterate 3000 rounds to find the correct answer with 99% prob
94     # assuming 20% outliers
95     optimal_H = None
96     max_inliers = 0
97     for i in range(3000):
98         # select 3 matches at random to calculate homography
99         matches = random.sample(matched_keypoints, 3)
100         H = calculate_homography(matches)
101
102         # count how many matches are inliers
103         num_inliers = count_inliers(H, matched_keypoints)
104
105         # get the homography with the most inliers
106         if num_inliers > max_inliers:
107             max_inliers = num_inliers
108             optimal_H = H
109
110     return optimal_H
111
112
113 def stitch_images(H, img1, img2):
114     """
115     Stitch two images together using the best homography H.
116     """
117     nrow_1, ncol_1 = img1.shape[0], img1.shape[1]
118     nrow_2, ncol_2 = img2.shape[0], img2.shape[1]
119
120     # transform 4 corners of the first image
121     (topleft_x, topleft_y) = transform_point(H, (0, 0))
122     (topright_x, topright_y) = transform_point(H, (ncol_1-1, 0))
123     (bottomright_x, bottomright_y) = transform_point(H, (ncol_1-1, nrow_1-1))
124     (bottomleft_x, bottomleft_y) = transform_point(H, (0, nrow_1-1))
125
126     # calculate dimensions

```

```

126     # calculate dimensions
127     min_y = min(topleft_y, topright_y, 0)
128     max_y = max(bottomleft_y, bottomright_y, nrow_2-1)
129     min_x = min(topleft_x, bottomleft_x, 0)
130     max_x = max(topright_x, bottomright_x, ncol_2-1)
131
132     # stitched image template
133     nrows_stitched = max_y - min_y + 1
134     ncols_stitched = max_x - min_x + 1
135     shape = (int(nrows_stitched)+1, int(ncols_stitched)+1, 3)
136     stitched_image = np.zeros(shape)
137
138     # compute offset
139     offset_x = 0 - min_x
140     offset_y = 0 - min_y
141
142     # transform pixels in the first image
143     for x in range(ncol_1):
144         for y in range(nrow_1):
145             (x_t, y_t) = transform_point(H, (x,y))
146             x_t = int(round(x_t + offset_x))
147             y_t = int(round(y_t + offset_y))
148             stitched_image[y_t, x_t, :] = img1[y, x, :]
149
150     # shift pixels in the second image
151     for x in range(ncol_2):
152         for y in range(nrow_2):
153             x_t = int(round(x + offset_x))
154             y_t = int(round(y + offset_y))
155             stitched_image[y_t, x_t, :] = img2[y, x, :]
156
157     # interpolate skipped pixels
158     for x in range(1, stitched_image.shape[1]-1):
159         for y in range(1, stitched_image.shape[0]-1):
160             # interpolate based on neighbours if skipped
161             if np.all(stitched_image[y, x, :] == 0):
162                 interpolate_x = 0.5*(stitched_image[y, x-1, :] +
163                                     stitched_image[y, x+1, :])
164                 interpolate_y = 0.5*(stitched_image[y-1, x, :] +
165                                     stitched_image[y+1, x, :])
166                 if np.linalg.norm(interpolate_x) > np.linalg.norm(interpolate_y):
167                     stitched_image[y, x, :] = interpolate_x
168                 else:
169                     stitched_image[y, x, :] = interpolate_y
170
171     return stitched_image
172
173
174 if __name__ == '__main__':
175     landscape1_path = 'landscape_1.jpg'
176     landscape2_path = 'landscape_2.jpg'
177     img1 = cv2.imread(landscape1_path)
178     img2 = cv2.imread(landscape2_path)
179
180     matched_keypoints = get_matching_keypoints(img1, img2)
181     H = RANSAC_homography(matched_keypoints)
182     stitch = stitch_images(H, img1, img2)
183
184     cv2.imwrite('stitched_output.jpg', stitch)

```

3. image um_000038.png


Q3.

(a) focal length $f = 721.5 \text{ mm} = 0.7215 \text{ m}$

principal point $(p_x = 0.6096 \text{ m}, p_y = 0.1729 \text{ m})$

→ internal camera parameter matrix

$$K = \begin{bmatrix} f & 0 & p_x \\ 0 & f & p_y \\ 0 & 0 & 1 \end{bmatrix} = \begin{bmatrix} 0.7215 & 0 & 0.6096 \\ 0 & 0.7215 & 0.1729 \\ 0 & 0 & 1 \end{bmatrix}$$



Thus, the equation of the ground plane in camera's coordinate system is:

$$Y = -1.7$$

(c)

$$\begin{bmatrix} w \cdot x \\ w \cdot y \\ w \end{bmatrix} = K \begin{bmatrix} X \\ Y \\ Z \end{bmatrix} \longrightarrow \begin{bmatrix} X \\ Y \\ Z \end{bmatrix} = w K^{-1} \begin{bmatrix} w \cdot x \\ w \cdot y \\ w \end{bmatrix} = \begin{bmatrix} 0.7215 & 0 & 0.6096 \\ 0 & 0.7215 & 0.1729 \\ 0 & 0 & 1 \end{bmatrix}^{-1} \begin{bmatrix} w \cdot x \\ w \cdot y \\ w \end{bmatrix}$$

$$= \begin{bmatrix} 1.386 & 0 & -0.845 \\ 0 & 1.386 & -0.24 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} w \cdot x \\ w \cdot y \\ w \end{bmatrix}$$

By assuming that the point lies on the ground, $Y = -1.7 = 1.386 w y - 0.24 w \longrightarrow w = \frac{-1.7}{1.386 y - 0.24}$

Thus, given 2D point (x, y) , the 3D location would be:

$$\begin{bmatrix} X \\ Y \\ Z \end{bmatrix} = \begin{bmatrix} 0.7215 & 0 & 0.6096 \\ 0 & 0.7215 & 0.1729 \\ 0 & 0 & 1 \end{bmatrix}^{-1} \begin{bmatrix} w \cdot x \\ w \cdot y \\ w \end{bmatrix} = \begin{bmatrix} 1.386 & 0 & -0.845 \\ 0 & 1.386 & -0.24 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} w \cdot x \\ w \cdot y \\ w \end{bmatrix}; \text{ where } w = \frac{-1.7}{1.386 y - 0.24}$$

4. Describe (in mathematical form, no code) how to compute disparity for a pair of parallel stereo cameras.

- Point P in the world P (X, Y, Z)
- O_l : left camera center
- O_r : right camera center
- $P_l(x_l, y_l)$: point on the left image ($w_1 \times h_1$)
- $P_r(x_r, y_r)$: point on the right image ($w_2 \times h_2$)
- Since lines O_lO_r and P_lP_r are parallel, and O_l and O_r have the same y, then also P_l and P_r have the same y, that is, $y_l = y_r$.
- Similar triangles: $\frac{T}{Z} = \frac{T+x_l-x_r}{Z-f}$; depth and disparity are inversely proportional
- `Compute_disparity(left_img, right_img)`:
 - For each point $P_l(x_l, y_l)$ in the `left_img`:
 - Find its corresponding point $P_r(x_r, y_r)$ in the `right_img` matching on the scanline $y_l = y_r$ by using SSD or normalized correlation
 - Calculate disparity $x_r - x_l$
 - Calculate depth $Z = \frac{f \times T}{x_r - x_l}$, where $T = O_r - O_l$
- Computation Complexity:
 - $SSD(patch_l, patch_r) = \sum_x \sum_y \left(I_{patch_l}(x, y) - I_{patch_r}(x, y) \right)^2$
 - $NC(patch_l, patch_r) = \frac{\sum_x \sum_y (I_{patch_l}(x, y) \times I_{patch_r}(x, y))}{|I_{patch_l}| \times |I_{patch_r}|}$
 - $O(w_1 w_2 h_1 h_2)$